

Class08: Breast Cancer Mini Project

Qiman A17912409

Background

The goal of this mini-project is for you to explore a complete analysis using the unsupervised learning techniques covered in class.

Today we will analyze a biopsy data-set from the needle aspiration (FNA)

Data import

```
# Save your input data file into your Project directory
fna.data <- "WisconsinCancer.csv"

# Complete the following code to input the data and store as wisc.df
wisc.df <- read.csv(fna.data, row.names=1)
```

Examine the input data to ensure column names are set correctly

```
head(wisc.df,4)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean
842302	0.11840	0.27760	0.3001	0.14710
842517	0.08474	0.07864	0.0869	0.07017
84300903	0.10960	0.15990	0.1974	0.12790
84348301	0.14250	0.28390	0.2414	0.10520

	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
--	---------------	------------------------	-----------	------------	--------------

842302	0.2419		0.07871	1.0950	0.9053	8.589
842517	0.1812		0.05667	0.5435	0.7339	3.398
84300903	0.2069		0.05999	0.7456	0.7869	4.585
84348301	0.2597		0.09744	0.4956	1.1560	3.445
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	153.40	0.006399	0.04904	0.05373		0.01587
842517	74.08	0.005225	0.01308	0.01860		0.01340
84300903	94.03	0.006150	0.04006	0.03832		0.02058
84348301	27.23	0.009110	0.07458	0.05661		0.01867
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst		
842302	0.03003		0.006193	25.38		17.33
842517	0.01389		0.003532	24.99		23.41
84300903	0.02250		0.004571	23.57		25.53
84348301	0.05963		0.009208	14.91		26.50
	perimeter_worst	area_worst	smoothness_worst	compactness_worst		
842302	184.60	2019.0		0.1622		0.6656
842517	158.80	1956.0		0.1238		0.1866
84300903	152.50	1709.0		0.1444		0.4245
84348301	98.87	567.7		0.2098		0.8663
	concavity_worst	concave.points_worst	symmetry_worst			
842302	0.7119		0.2654			0.4601
842517	0.2416		0.1860			0.2750
84300903	0.4504		0.2430			0.3613
84348301	0.6869		0.2575			0.6638
	fractal_dimension_worst					
842302		0.11890				
842517		0.08902				
84300903		0.08758				
84348301		0.17300				

Make sure we don't accidentally include diagnosis column in our analysis.

```
# We can use -1 here to remove the first column
wisc.data <- wisc.df[,-1]
```

Now we set up a new vector called diagnosis and store this as a factor for later.

```
# Create diagnosis vector for later
diagnosis <- as.factor(wisc.df$diagnosis)
```

Exploratory Data Analysis

Q1.How many observations are in this dataset?

```
nrow(wisc.data)
```

```
[1] 569
```

There are 569 observations in this dataset.

Q2.How many of the observations have a malignant diagnosis?

```
table(wisc.df$diagnosis)
```

```
  B   M  
357 212
```

There are 212 observations with a malignant diagnosis.

Q3. How many variables/features in the data are suffixed with `_mean`?

We can use the `grep()` function to help us here:

```
length(grep("_mean",colnames(wisc.data), value=TRUE))
```

```
[1] 10
```

Principal Component Analysis (PCA)

We need to scale our data before PCA with the `scale=TRUE` argument to `prcomp()`.

```
wisc.pr <- prcomp(wisc.data, scale=TRUE)  
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)?

The proportion of the original variance is .4427 so 44.27% of the total variance in the dataset is captured by PC1.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?


You need 3 principal components to describe at least 70% of the original variance in the data.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

You need 7 principal components to describe at least 90% of the original variance in the data.

Creating a biplot:

```
biplot(wisc.pr)
```




Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

This plot is a huge mess and it is very difficult to understand. All the values are highly clustered and it is impossible to read the plot properly.

Let's see the "PC score plot"


```
library(ggplot2)

ggplot(wisc.pr$x)+
  aes(PC1,PC2, col=diagnosis)+
  geom_point()
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
ggplot(wisc.pr$x)+  
  aes(PC1,PC3, col=diagnosis)+  
  geom_point()
```



Overall the plots show that principal component 1 is capturing a separation of malignant and benign samples.

Variance Explained


Use a scree plot to show how much variance each PC captures. Look for an “elbow”.

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```


```
[1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

```
#Variance explained by each principal component: pve
pve <- pr.var/sum(pr.var)

# Plot variance explained for each principal component
plot(c(1,pve), xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



```
# Alternative scree plot of the same data, note data driven y-axis
barplot(pve, ylab = "Percent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```

Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

```
# Loading for concave.points_mean on PC1
wisc.pr$rotation["concave.points_mean", 1]
```

```
[1] -0.2608538
```

```
# Check largest absolute loadings for PC1
sort(abs(wisc.pr$rotation[,1]), decreasing = TRUE)
```

concave.points_mean	concavity_mean	concave.points_worst
0.26085376	0.25840048	0.25088597
compactness_mean	perimeter_worst	concavity_worst
0.23928535	0.23663968	0.22876753
radius_worst	perimeter_mean	area_worst
0.22799663	0.22753729	0.22487053


area_mean	radius_mean	perimeter_se
0.22099499	0.21890244	0.21132592
compactness_worst	radius_se	area_se
0.21009588	0.20597878	0.20286964
concave.points_se	compactness_se	concavity_se
0.18341740	0.17039345	0.15358979
smoothness_mean	symmetry_mean	fractal_dimension_worst
0.14258969	0.13816696	0.13178394
smoothness_worst	symmetry_worst	texture_worst
0.12795256	0.12290456	0.10446933
texture_mean	fractal_dimension_se	fractal_dimension_mean
0.10372458	0.10256832	0.06436335
symmetry_se	texture_se	smoothness_se
0.04249842	0.01742803	0.01453145

For the first principal component the loading for concave.points_mean is approximately -0.2609 meaning that concave.points_mean contributes strongly to PC1. There are no features with larger contributions than this one.

Hierarchical Clustering

```
wisc.hclust <- hclust( dist( scale(wisc.data)))
plot(wisc.hclust)
```

Cluster Dendrogram




```
dist(scale(wisc.data))  
hclust (*, "complete")
```

Q10. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)  
abline(h=18, col="red", lty=2)
```

Cluster Dendrogram




```
dist(scale(wisc.data))  
hclust (*, "complete")
```

at h=18 the clustering model has 4 clusters. ##Combining methods

```
d <- dist(wisc.pr$x[,1:4])  
wisc.pr.hclust <- hclust(d,method="ward.D2")  
plot(wisc.pr.hclust)  
abline(h=70, col="red")
```

Cluster Dendrogram



d
hclust (*, "ward.D2")

```
wisc.hclust.clusters <- cutree(wisc.hclust, h=20)
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

Q12. Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning.

ward.D2 gives the best results because it produce the clearest and most interpretable cluster structure. Te clusters were more compact and distinct.

```
grps <- cutree(wisc.pr.hclust, h=70)
table(grps)
```


grps	1	2
	171	398

How does this clustering `grps` correspond to the expert diagnosis

```
table(diagnosis, grps)
```

```
      grps
diagnosis  1  2
B      6 351
M     165  47
```

```
ggplot(wisc.pr$x) +
  aes(PC1, PC2) +
  geom_point(col=grps)
```



```
l <- dist(wisc.pr$x[,1:7])
wisc.pr.hclust1 <- hclust(l,method="ward.D2")
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust1, k=2)
```

Q13. How well does the newly created `hclust` model with two clusters separate out the two “M” and “B” diagnoses?

```
table(wisc.pr.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.pr.hclust.clusters	B	M
1	28	188
2	329	24

The two cluster model separates the diagnoses well but not perfectly. Cluster 1 is mostly malignant with 28 benign. Cluster two is mostly benign with 24 malignant. The model does distinguish between the two groups fairly well with the accuracy being high.

Q14. How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.hclust.clusters` and `wisc.pr.hclust.clusters`) with the vector containing the actual diagnoses.

```
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

```
table(wisc.pr.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.pr.hclust.clusters	B	M
1	28	188
2	329	24

The hierarchical clustering model without PCA can separate benign and malignant diagnoses but the clusters are more fragmented and less interpretable. For example small clusters 2 and 4 contain few samples, there is some separation but it is fragmented into 4 clusters. The PCA-based clustering performs better overall giving a clearer separation between the two groups. One cluster is mostly malignant while the other is mostly benign. This is a much cleaner 2-group separation.


Prediction

We will use the `predict()` function that will take our PCA model from before and new cancer cell data and project that data onto our PCA space.

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	
	PC27	PC28	PC29	PC30			
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029			
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820			

```
plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```

Q16. Which of these new patients should we prioritize for follow up based on your results?

We should prioritize patient 2 because it has $PC1 = -4.25$ which places it closer to the cluster of malignant samples. As seen in the plot patient 2 clusters closer to the malignant tumor samples.